

Winter term 2025/2026

Practical parallel algorithms with MPI

Dr. Jan Eitzinger

Erlangen National High Performance Computing Center (NHR@FAU)

Lecture 6: Advanced MPI: derived data types and virtual topologies



Motivation heterogenous data

Example: Root reads configuration and broadcasts it to all others

```
// root: read configuration from  
// file into struct config  
  
MPI_Bcast(&cfg.nx, 1, MPI_INT, ...);  
MPI_Bcast(&cfg.ny, 1, MPI_INT, ...);  
MPI_Bcast(&cfg.du, 1, MPI_DOUBLE,...);  
MPI_Bcast(&cfg.it, 1, MPI_INT, ...);
```



Want to do something like:

```
MPI_Bcast(  
    &cfg, 1, <type cfg>, ...);
```

```
        MPI_Bcast(&cfg, sizeof(cfg),  
                  MPI_BYTE, ...)
```

is **not** a solution. It's not portable as no
data conversion can take place

Motivation non-contiguous data

- Example: Send column of matrix (noncontiguous in C):
 - Send each element alone?
 - Manually copy elements out into a contiguous buffer and send it?

0	1	2	3	4
5	6	7	8	9
10	11	12	13	14
15	16	17	18	19
20	21	22	23	24
25	26	27	28	29

MPI general datatypes

- Datatypes allow to (de)serialize arbitrary data layouts into a message stream
 - Networks provide serial channels
 - Same for block devices and I/O
- Powerful, generic specification to describe arbitrary data layouts
- Declarative specification of data-layout “what” and not “how”, leaves optimization to implementation
- Derived datatypes constructed from basic datatypes using constructors
- Recursive concept: Derived data types can be constructed from other derived datatypes

Specification

- A general MPI datatype object consists of a pair of
 - Sequence of basic datatypes
 - Sequence of integer (byte) displacements
- Nomenclature:
 - Type map: Sequence of pairs (datatype, displacement)

```
{(int, 64), (double, 8), (double, 8), (double, 48)};
```

- Type signature: Sequence of datatypes (displacements ignored)

```
{int, double, double, double};
```

Need not be
distinct, positive or
in increasing order!


Example

Typemap datatypeA

```
{(int, 0), (double, 8), (int, 64)};
```

Typemap datatypeB

```
{(int, 68), (double, 0), (int, 64)};
```



MPI datatype extent

- First byte to last byte occupied, rounded to alignment requirements

$$\text{Typemap} = \{ (\text{type}_0, \text{disp}_0), \dots, (\text{type}_{n-1}, \text{disp}_{n-1}) \},$$
$$\text{lb}(\text{Typemap}) = \min(\text{disp}_j)$$
$$\text{ub}(\text{Typemap}) = \max(\text{disp}_j + \text{sizeof}(\text{type}_j)) + \epsilon$$
$$\text{extent}(\text{Typemap}) = \text{ub}(\text{Typemap}) - \text{lb}(\text{Typemap})$$

Example

```
{ (double, 0), (char, 8) };
```

Assumption: doubles
are 8-byte aligned!

Extent of datatype is 16.

Creating an MPI data type

Three steps:

1. Construct with

```
MPI_Type_*(...);
```

2. Commit new data type with

```
MPI_Type_commit(MPI_Datatype * nt);
```


3. After use, deallocate the data type with

```
MPI_Type_free(MPI_Datatype * nt);
```

All local, non-collective calls

The simplest type: MPI_Type_contiguous

```
MPI_Type_contiguous(int count, MPI_Datatype oldtype,  
                 MPI_Datatype * newtype);
```



Can also be
general derived
MPI type!

```
MPI_Datatype nt;  
MPI_Type_contiguous(  
  3, MPI_INT, &nt);  
  
MPI_Type_commit(&nt);  
// use nt...  
MPI_Type_free(&nt);
```

A flexible, vector-like type: MPI_Type_vector

```
MPI_Type_vector(int count, int blocklength, int stride,  
                MPI_Datatype oldtype,  
                MPI_Datatype * newtype);
```


count 2 (no. of blocks)

blocklength 3 (no. of elements in each block)

stride 5 (no. of elements b/w start of each block)

oldtype


MPI_INT



```
MPI_Datatype nt;  
MPI_Type_vector(  
 2, 3, 5, MPI_INT, &nt);  
  
MPI_Type_commit(&nt);  
// use nt...  
MPI_Type_free(&nt);
```

Sending a column of a matrix in C

Row-major data layout in C → cannot use plain array



```
double matrix[30];
MPI_Datatype nt;

// count = nrows, blocklength = 1,
// stride = ncols
MPI_Type_vector(nrows, 1, ncols,
                 MPI_DOUBLE, &nt);
MPI_Type_commit(&nt);

// send column
MPI_Send(&matrix[1], 1, nt, ...);

MPI_Type_free(&nt);
```

An even more flexible type: **`MPI_Type_indexed`**

```
MPI_Type_indexed(int count,
int blocklengths[], int displacements[],
MPI_Datatype oldtype, MPI_Datatype * newtype);
```

<code>int count</code>	number of blocks, also number of entries in array of displacements and array of blocklengths
<code>int blocklengths[]</code>	number of elements per block (array of non-negative integers)
<code>int displacements[]</code>	displacement for each block, in multiples of oldtype (array of integers)

Simple example MPI_Type_indexed

```
Old typemap : { (double, 0), (double,8) };
```



```
extent : 16
```

```
count : 2
```

```
blocklengths : (3, 1)
```

```
displacements : (4, 0)
```

```
New typemap : { (double, 64), (double,72),  
                (double,80), (double,88), (double,96), (double,104),  
                (double,0), (double,8) };
```



A more realistic example


```
double a[100], b[100];
int displs[10];
int lens[10];
MPI_Datatype tril;
MPI_Status status;

for (int i = 0; i < 10; i++) {
    displs[i] = i * 10;
    lens[i]   = i + 1;
}

MPI_Type_indexed(10, lens, displs, MPI_DOUBLE, &tril);
MPI_Type_commit(&tril);

MPI_Sendrecv(a, 1, tril, rank, 0, b, 1, tril, rank, 0,
            MPI_COMM_WORLD, &status);
```

Copy the lower triangular part of a matrix between two arrays on same process



Constructor variants using addresses

- Type constructors provide variant with displacement in bytes instead of element extent:
 - `MPI_Type_create_hvector`
 - `MPI_Type_create_hindexed`
- Special MPI type for addresses and displacements: `MPI_Aint`
- Always use MPI provided routines for address computations!
- Usually, addresses are relative to buffer address provided to communication routines
- It is allowed to use absolute addresses, in this case set buffer address = `MPI_BOTTOM`

How to obtain and handle addresses?

```
MPI_Get_address(const void *location, MPI_Aint *address);  
MPI_Aint MPI_Aint_diff(MPI_Aint addr1, MPI_Aint addr2);  
MPI_Aint MPI_Aint_add(MPI_Aint base, MPI_Aint disp);
```

- Example:

```
double a[100];  
MPI_Aint a1, a2, disp;  
MPI_Get_address(&a[0], &a1);  
MPI_Get_address(&a[50], &a2);  
disp = MPI_Aint_diff(a2,a1);
```

Result would usually be `disp = 400 (50 x 8)`

Derived type size and extent


- Get the total **size** (in bytes) of datatype in a message (Type signature)

```
int MPI_type_size(MPI_Datatype newtype, int *size);
```

- Get the lower bound and the **extent** (span from the first byte to the last byte) of datatype

```
int MPI_type_get_extent(MPI_Datatype newtype,  
                        MPI type for  
                        memory addresses  
                        or offsets {  
                           MPI_Aint *lb,  
                           MPI_Aint *extent);
```


- MPI allows to change the **extent** of a datatype
 - using **lb_marker** and **ub_marker**
 - does not affect the size or count of a datatype, and the message content
 - does affect the outcome of a replication of this datatype



Most flexible type: MPI_Type_create_struct

Describe blocks with arbitrary data types and arbitrary displacements

```
MPI_Type_create_struct(int count, int block_lengths[],  
MPI_Aint displs[], MPI_Datatype types[],  
MPI_Datatype * newtype);
```



The contents of `displs` are either the displacements in bytes of the block bases or MPI addresses


A sub-array type: `MPI_Type_create_subarray`

```
MPI_Type_create_subarray(int dims,
    int ar_sizes[], int ar_subsizes[], int ar_starts[],
    int order, MPI_Datatype oldtype, MPI_Datatype * newtype);
```

- **dims**: dimension of the array
- **ar_sizes**: array with sizes of array (dims entries)
- **ar_subsizes**: array with sizes of subarray (dims entries)
- **ar_starts**: start indices of the subarray inside array (dims entries), start at 0 (also in Fortran)
- **order**
 - row-major: `MPI_ORDER_C`
 - column-major: `MPI_ORDER_FORTRAN`

Example for a sub-array type: “bulk” of a matrix

dims	2
ar_sizes	{ncols, nrows}
ar_subsizes	{ncols-2, nrows-2}
ar_starts	{1, 1}
order	MPI_ORDER_C
oldtype	MPI_INT



```
MPI_Type_create_subarray(dims, ar_sizes, ar_subsizes,  
                        ar_starts, order, oldtype, &nt);  
MPI_Type_commit(&nt);  
// use nt...  
MPI_Send(&buf[0], 1, nt, ...); // etc.  
MPI_Type_free(&nt);
```

Performance considerations

- Derived datatypes enable **zero copy**: MPI forces “remote” copy
- MPI implementations **copy internally**
 - E.g., networking stack (TCP), packing derived data types
 - **Zero-copy is possible** (RDMA, I/O Vectors, SHMEM)
- MPI applications **copy too often**
 - E.g., manual pack, unpack or data rearrangement
 - **Derived datatypes can do both!**
- **Simple and effective performance model:**
 - **More parameters == slower**

Some implementations
might still be slower than
manual packing/unpacking!


Virtual (Cartesian) topologies in MPI

A convenient process naming scheme for multi-dimensional problems

- Convenient process naming
- Naming scheme to fit the communication pattern
- Simplifies writing of code
- Can allow MPI to optimize communications

- Let MPI map ranks to coordinates
- User: map array segments to ranks

Example: distribute
2-D array of
4000 x 3000 elements
equally on 12 ranks



Creating a new Cartesian communicator

- Create new communicator attached to Cartesian topology

```
MPI_Cart_create(MPI_Comm oldcomm,  
                 int ndims, int dims[], int periods[],  
                 int reorder, MPI_Comm *cart_comm);
```

ndims: number of dimensions

dims: array with **ndims** elements,

dims[i] specifies the number of ranks in dimension **i**

periods: array with **ndims** elements,


periods[i] specifies if dimension **i** is periodic

reorder: allow rank of **oldcomm** to have a different rank in **cart_comm**

Cartesian topology example

```
ndims      = 2  
dims       = {4, 3}  
periods   = {1, 0}  
reorder    = 0
```

Coordinates



Convenience Function: `MPI_Dims_create`

- Select a balanced distribution of processes per coordinate direction

```
MPI_Dims_create(
    int nnodes, int ndims, int dims[]);
```

`nnodes` : number of nodes in grid, usually the number of processes
`ndims` : number of cartesian dimensions
`dims` : array with `ndims` elements, number of nodes in each dimension

Argument `dims` allows to control the layout:

- If `dims[i]` is set to a positive number, # nodes in dimension i is not touched
- Only entries where `dims[i]` is set to 0 are modified by call! There is an **error** if `dims[i]` is not initialized!

Cartesian topology service functions

- Retrieve rank in new Cartesian communicator (“who am I in the grid?”)

```
MPI_Comm_rank(cart_comm, int *cart_rank);
```

- Map rank → coordinates (“where am I in the grid? ”)

```
MPI_Cart_coords(cart_comm, rank, int maxdims, int coords[]);
```

rank: any rank which is part of Cartesian communicator **cart_comm**

coords: array of **maxdims** elements, receives the coordinates for **rank**


- Map coordinates → rank (“who is at that position? ”)

```
MPI_Cart_rank(cart_comm, int coords[], int *rank);
```

coords: coordinates; if periodic in direction **i**, **coords[i]** are automatically mapped into the valid range, else they are erroneous

Example

- Example: 12 processes arranged on a 4×3 grid
- Process coordinates begin with 0



Next-neighbor communication

Sending/receiving from neighbors is a typical task in Cartesian topologies

```
MPI_Cart_shift(cart_comm, direction, disp,  
                  int *source_rank, int *dest_rank);
```

direction: dimension to shift

disp: offset to shift: > 0 shift in positive direction,
< 0 shift in negative direction

source_rank/dest_rank: returned ranks as input
into **MPI_Sendrecv*** calls

Next-neighbor communication

Example: 4x3 process grid, periodic in 1st dimension, each process has an `int` value, which gets shifted

```
MPI_Cart_shift(cart_comm, 0, 1, &src, &dst);  
MPI_Sendrecv_replace(&value, 1, MPI_INT,  
                     dst, 0, src, 0, cart_comm, ...)
```

9	0	3	6	9	0
10	1	4	7	10	1
11	2	5	8	11	2

Rank 0:
src: 9
dst: 3

```
MPI_Cart_shift(cart_comm, 1, 1, &src, &dst);  
MPI_Sendrecv_replace(&value, 1, MPI_INT,  
                     dst, 0, src, 0, cart_comm, ...)
```

-2	-2	-2	-2
0	3	6	9
1	4	7	10
2	5	8	11
-2	-2	-2	-2

Rank 0:
src: 3
dst: MPI_PROC_NULL

for non-periodic dimensions
`MPI_PROC_NULL` is
returned on boundaries

shift in 1st dimension, which is periodic

shift in 2nd dimension, which is non-periodic

Putting it all together

Setup virtual topology for cartesian 2D grid:

```
int dims[2] = {0, 0};  
int periods[2] = {0, 0};  
MPI_Dims_create(size, 2, dims);  
MPI_Cart_create(MPI_COMM_WORLD, 2, dims, periods, 0, &comm);  
MPI_Cart_shift(comm, 0, 1, &iNeighbours[0], &iNeighbours[1]);  
MPI_Cart_shift(comm, 1, 1, &jNeighbours[0], &jNeighbours[1]);  
MPI_Cart_get(comm, 2, dims, periods, coords);
```

Check if rank is at specific boundary:

```
if ( coords[1] == (dims[1]-1) ) { //set top bc  
...  
}
```

Derived MPI datatypes for ghost layers

- Compute size of local domain


```
imaxLocal = sizeOfRank(rank, dims[0], imax);  
jmaxLocal = sizeOfRank(rank, dims[1], jmax);
```

- Create type for contiguous dimension i

```
MPI_Type_contiguous(imaxLocal, MPI_DOUBLE,  
                    &jBufferType);  
MPI_Type_commit(&jBufferType);
```

- Create type for non-contiguous dimension j

```
MPI_Type_vector(jmaxLocal, 1, imaxLocal+2,  
                  MPI_DOUBLE, &iBufferType);  
MPI_Type_commit(&iBufferType);
```



Example: Exchange boundary

```
double* buf[8]; MPI_Request requests[8];
for ( int i=0; i<8; i++ ) requests[i] = MPI_REQUEST_NULL;
buf[0] = grid + 1; //recv bottom
buf[1] = grid + (imaxLocal+2) + 1; //send bottom
// Setup send and receive buffers

for (int i=0; i<2; i++) {
    /* exchange ghost cells with bottom/top neighbor */
    MPI_Irecv(buf[i*2], 1, jBufferType, jNeighbours[i], 1,
              comm, &requests[i*2]);
    MPI_Isend(buf[(i*2)+1], 1, jBufferType, jNeighbours[i], 1,
              comm, &requests[i*2+1]);
    /* exchange ghost cells with left/right neighbor */
    MPI_Irecv(buf[i*2+4], 1, iBufferType, iNeighbours[i], 1,
              comm, &requests[i*2+4]);
    MPI_Isend(buf[i*2+5], 1, iBufferType, iNeighbours[i], 1,
              comm, &requests[(i*2)+5]);
}
MPI_Waitall(8, requests, MPI_STATUSES_IGNORE);
```

What else? Covered later in lecture!


Support for general graph topologies

- Constructors `MPI_Graph_create` and `MPI_Dist_graph_create`
- `MPI_TOPO_TEST` is used to query for the type of topology associated with a communicator
- `MPI_GRAPHDIMS_GET` and `MPI_GRAPH_GET` to retrieve graph topology information
- `MPI_GRAPH_NEIGHBORS_COUNT` and `MPI_GRAPH_NEIGHBORS` to obtain the neighbors of an arbitrary node in the graph
- `MPI_DIST_GRAPH_NEIGHBORS_COUNT` and `MPI_DIST_GRAPH_NEIGHBORS` to obtain the neighbors of the calling process

Neighborhood Collective Communication on Process Topologies

MPI-3 feature: Neighborhood collective

- Nearest neighbor communication that can be expressed as MPI virtual topologies
- Opportunity for optimized mapping on network topology
- Collective operations, must be called by all processes in communicator
- With Cartesian topology the sequence of neighbors in buffers is defined by
 - order of the dimensions
 - first the neighbor in the negative direction
 - then in the positive direction



Simplest Neighborhood collective: allgather

- For all direct neighbor ranks:

Send the same data to all ranks

Receive data from all ranks

```
MPI_Neighbor_allgather(  
    void* sendbuf, int sendcount, MPI_Datatype sendtype,  
    void* recvbuf, int recvcount, MPI_Datatype recvtype,  
    MPI_Comm comm);
```

sendbuf, sendcount, sendtype: Send data specification.


recvbuf, recvcount, recvtype: Receive data specification.

recvcount: number of elements received from each neighbor.

comm: communicator created with **MPI_Cart_create**


Neighborhood collective variations allgather

MPI_Neighbor_allgather



Contiguous chunks
from each neighbor


MPI_Neighbor_allgatherv




Non-contiguous variable-sized
chunks from each neighbor

Neighborhood collective variations alltoall

`MPI_Neighbor_alltoall`



`MPI_Neighbor_alltoallv`



Contiguous chunks in send and recv buffers for each neighbor

Non-contiguous variable-sized chunks in send and recv buffers for each neighbor

Neighborhood collective **`MPI_Neighbor_alltoallw`**

```
MPI_Neighbor_alltoallw(
    void* sendbuf, int sendcounts[], MPI_Aint sdispls[],
    MPI_Datatype sendtypes[],
    void* recvbuf, int recvcounts[], MPI_Aint rdispls[],
    MPI_Datatype recvtypes[],
    MPI_Comm comm);
```

`sendbuf, sendcounts, sdispls, sendtypes`: Send data specification.

`recvbuf, recvcounts, rdispls, recvtypes`: Receive data specification

`comm`: communicator created with **`MPI_Cart_create`**

In bytes!

2D Halo exchange with MPI_Neighbor_alltoallw


- Communication setup

```
int counts[4] = {1, 1, 1, 1};

bufferTypes[0] = vectorType;           //left
bufferTypes[1] = vectorType;           //right
bufferTypes[2] = contiguousType;        //bottom
bufferTypes[3] = contiguousType;        //top

sdispls[0] = ((imaxLocal+2)+1)*dblsize;
sdispls[1] = ((imaxLocal+2)+imaxLocal)*dblsize;
sdispls[2] = ((imaxLocal+2)+1)*dblsize;
sdispls[3] = ((jmaxLocal)*(imaxLocal+2)+1)*dblsize;


rdispls[0] = (imaxLocal+2)*dblsize;
rdispls[1] = ((imaxLocal+2)+(imaxLocal+1))*dblsize;
rdispls[2] = 1*dblsize;
rdispls[3] = ((jmaxLocal+1)*(imaxLocal+2)+1))*dblsize;
```



2D Halo exchange with MPI_Neighbor_alltoallw

- Communication exchange

```
MPI_Neighbor_alltoallw(  
buf, counts, sdispls, bufferTypes,  
buf, counts, rdispls, bufferTypes, comm );
```



Simple debugging of communication

With small domain size:

1. Initialize array with rank id
2. Perform exchange once
3. Print array

```
for ( int i=0; i<(imaxLocal+2)*(jmaxLocal+2) ; i++ ) {  
    buf[i] = rank;  
}  
  
exchange(buf);  
print(buf);
```

```
20: 1.00000000 1.00000000 1.00000000 1.00000000 1.00000000 1.00000000 1.00000000  
21: 1.00000000 1.00000000 1.00000000 1.00000000 1.00000000 1.00000000 1.00000000  
### RANK 2 #####  
00: 2.00000000 2.00000000 2.00000000 2.00000000 2.00000000 2.00000000 2.00000000  
01: 0.00000000 2.00000000 2.00000000 2.00000000 2.00000000 2.00000000 2.00000000  
02: 0.00000000 2.00000000 2.00000000 2.00000000 2.00000000 2.00000000 2.00000000  
03: 0.00000000 2.00000000 2.00000000 2.00000000 2.00000000 2.00000000 2.00000000  
04: 0.00000000 2.00000000 2.00000000 2.00000000 2.00000000 2.00000000 2.00000000  
05: 0.00000000 2.00000000 2.00000000 2.00000000 2.00000000 2.00000000 2.00000000  
06: 0.00000000 2.00000000 2.00000000 2.00000000 2.00000000 2.00000000 2.00000000  
07: 0.00000000 2.00000000 2.00000000 2.00000000 2.00000000 2.00000000 2.00000000  
08: 0.00000000 2.00000000 2.00000000 2.00000000 2.00000000 2.00000000 2.00000000  
09: 0.00000000 2.00000000 2.00000000 2.00000000 2.00000000 2.00000000 2.00000000
```

```
20: 0.00000000 2.00000000 2.00000000 2.00000000 2.00000000 2.00000000 2.00000000  
21: 2.00000000 3.00000000 3.00000000 3.00000000 3.00000000 3.00000000 3.00000000  
### RANK 3 #####  
00: 3.00000000 2.00000000 2.00000000 2.00000000 2.00000000 2.00000000 2.00000000  
01: 1.00000000 3.00000000 3.00000000 3.00000000 3.00000000 3.00000000 3.00000000  
02: 1.00000000 3.00000000 3.00000000 3.00000000 3.00000000 3.00000000 3.00000000  
03: 1.00000000 3.00000000 3.00000000 3.00000000 3.00000000 3.00000000 3.00000000  
04: 1.00000000 3.00000000 3.00000000 3.00000000 3.00000000 3.00000000 3.00000000  
05: 1.00000000 3.00000000 3.00000000 3.00000000 3.00000000 3.00000000 3.00000000  
06: 1.00000000 3.00000000 3.00000000 3.00000000 3.00000000 3.00000000 3.00000000
```

Debug output of distributed array

```
for ( int i=0; i < size; i++ ) {
    if ( i == rank ) {
        printf("### RANK %d #####\n", rank);
        for( int j=0; j < jmaxLocal+2; j++ ) {
            printf("%02d: ", j);
            for( int i=0; i < solver->imaxLocal+2; i++ ) {
                printf("%12.8f ", grid[j*(imaxLocal+2) + i]);
            }
            printf("\n");
        }
        fflush( stdout );
    }
    MPI_Barrier(MPI_COMM_WORLD);
}
```